

IBM 1Q 2026 Earnings

Conference Call Transcript

IBM Participants:

Olympia McNerney – Global Head of Investor Relations

Arvind Krishna – Chairman, President and Chief Executive Officer

James Kavanaugh – SVP and Chief Financial Officer

April 22, 2026 / 5:00 p.m. ET

MANAGEMENT DISCUSSION

Operator: Welcome and thank you for standing by. At this time, all participants are in a listen-only mode. Today's conference is being recorded. If you have any objections, you may disconnect at this time. Now, I will turn the meeting over to Olympia McNerney, IBM's Global Head of Investor Relations. Olympia, you may begin.

Olympia McNerney, IBM: Thank you. I'd like to welcome you to IBM's first quarter 2026 earnings presentation. I'm Olympia McNerney, and I'm here today with Arvind Krishna, IBM's Chairman, President and Chief Executive Officer; and Jim Kavanaugh, IBM's Senior Vice President and Chief Financial Officer.

We'll post today's prepared remarks and a replay of today's webcast on the IBM Investor website within a couple of hours. The earnings presentation is already available. To provide additional information to our investors, our presentation includes certain non-GAAP measures. For example, all of our references to revenue and signings growth are at constant currency. We provided reconciliation charts for these and other non-GAAP financial measures at the end of the presentation, which is posted to our Investor website.

Finally, some comments made in this presentation may be considered forward-looking under the Private Securities Litigation Reform Act of 1995. These statements involve factors that could cause our actual results to differ materially. Additional information about these factors is included in the company's SEC filings. So, with that, I'll turn the call over to Arvind.

Arvind Krishna, IBM: Thank you for joining us today. Let me start with our first quarter results and then provide context on what we are seeing across the business. IBM is off to a strong start to 2026, revenue in the first quarter grew 6% and combined with strong margin expansion, drove 13% growth in free cash flow. These results reflect the durability of our portfolio, the mission-critical nature of the work we do for clients, and the continued execution of our strategy.

Let me first touch on the macro. While we are operating in a dynamic environment, Middle East developments didn't impact us in the first quarter. Uncertainties remain, but our diversity across businesses, geographies, industries, and large enterprise clients position us well. Conversations we are having with clients remain consistent. Enterprises are investing in capabilities that increase resiliency, productivity and accelerate growth.

They're modernizing core systems. They are scaling AI, and they're making deliberate choices about where workloads should run and who controls the infrastructure underneath them. These are structural priorities, and they align directly with IBM's strengths. This quarter's performance reinforces the strategic choices we have made over the last several years to advance IBM as a software-led, hybrid cloud, and AI platform company.

Software revenue grew 8%, with Data and Red Hat growing double-digits. Infrastructure grew 12%, with another record Z quarter up 48%. We also had strong performance in Distributed Infrastructure as generative AI increases demand for our storage offerings. Consulting grew 1%, with momentum in enterprise data and business application transformations as clients modernize to deploy AI securely and at scale. The durability of our portfolio is a defining feature of IBM today.

Let me spend a few minutes on AI. Enterprises are still figuring out where to deploy this technology and where competitive advantage truly sits. Every major technology wave has followed a pattern. Value begins with infrastructure, moves to enabling platforms, and ultimately, concentrates in workflows where businesses operate. Right now, the spotlight is on foundation models.

Enterprises are building portfolios, frontier models for some workloads, smaller models running on-premises for others, and open-source models where control and flexibility matter the most. Enterprises will want to retain control of their proprietary data. AI will run everywhere – across public cloud, private and sovereign clouds, and on-premises.

The core challenge is making all of this work together. This includes orchestrating across models, agents, and workflows, governing enterprise data, and securing these systems at scale. And that is exactly where IBM operates. We are building the platform that lets enterprises put AI to work on their terms – wherever it runs, whichever models they choose, and under governance they control. Our portfolio is built around world-class security, support, and integration for an enterprise environment.

Red Hat provides a common open platform that lets enterprises run applications in AI consistently across any infrastructure. More AI adoption means more demand for open, flexible infrastructure. In automation, the logic is similar: agents multiply applications, integrations and execution paths. Managing that sprawl requires a control plane to provision infrastructure, integrate applications, secure environments, and manage cost. This is what our end-to-end automation portfolio provides.

In an AI-driven world, security risks are rising. IBM Concert identifies vulnerabilities proactively and automates remediation, helping enterprises maintain resilience at scale. Our data business is seeing similar AI tailwinds. AI is only as good as the data it can access, and increasingly, that data is not static. It is generated continuously across transactions, applications, and interactions. To deliver real-time AI outcomes, data must be available in motion, governed and delivered securely to models and agents wherever they are running.

Confluent, which we closed this past quarter, solves that directly. It streams live governed data to models and agents across the hybrid environment, and the orchestration layer ties it together. In a multi-model world, clients need to route between models, manage agent workflows, and maintain governance. That is what watsonx orchestrate and a watsonx platform deliver.

We have also created AI editions of critical software products like Db2, Cognos, and MQ. These embed agentic AI that can reason, act, and automate at scale, while preserving IBM-grade security and trust.

Infrastructure remains a critical differentiator, as AI moves into the core of enterprise operations. IBM Z delivers the lowest unit cost architecture at scale for workloads that require end-to-end encryption, continuous availability, and ultra-high throughput. Clients rely on our Z platform to process billions of transactions reliably, with 6 to 8 nines of availability. They run AI inferencing directly in line with those transactions.

Our Spyre Accelerator lets clients run AI on 100% of their transaction volume without moving data off platform, allowing them to embed AI directly into their transaction flows. Financial services clients are using this for real-time fraud detection, saving tens of millions of dollars. At the same time, AI assisted modernization, including code understanding, refactoring, and API integration, makes it easier to evolve applications without compromising the guarantees the platform provides.

Our watsonx Assistants for Z, were made available over two years ago, they help clients preserve the architectural strengths that deliver resilience, security, and scalability, while making the platform more productive. Clients who have deployed watsonx Code Assistant for Z are growing MIPS capacity three times faster than those who have not.

In Consulting, AI is both a growth driver and a productivity engine as agents take on more work, delivery becomes faster, more software-driven and more scalable. IBM is leaning into this shift through our Consulting Advantage platform and unique integrated value, sitting side-by-side with software. This helps clients operationalize AI while improving our own efficiency. Demand continues to accelerate as clients move beyond experimentation and focus on transforming applications, data, and workflows to embed AI into core operation.

All of this allows us to drive value for clients. ServiceNow is leveraging watsonx for automated data quality and observability to deliver AI-ready data and code generation to refresh legacy applications to modern application runtimes, including ServiceNow. Visa continues to work with IBM on ongoing software and data modernization efforts, supporting the scale, resiliency, and performance of VisaNet. With Nestlé, we are using NVIDIA accelerated watsonx.data to embed AI directly into core order to cash operations, enabling faster real-time insights across Nestlé's global supply chain. This highlights how quickly we can bring research to bear for commercial value. Nestlé was ideal for this proof of concept because of its strong digital backbone. In Infrastructure, clients such as NatWest and RBC are modernizing their mainframe environments, using AI and automation capabilities, including watsonx Assistant and watsonx Code Assistant for Z to improve resiliency, security and developer productivity.

We continue to accelerate organic innovation. IBM Bob, our AI-based software development system, is now generally available. Our entire developer workforce is using Bob with average productivity gains of 45%. Bob automates the full software life cycle from legacy modernization to security, using specialized agents and multi-modal optimization. It drives developer productivity and predictable costs.

We also introduced Sovereign Core, software that lets organizations run AI workloads under their own operational authority, within a defined jurisdiction and auditable controls. We see sovereignty

becoming a bigger factor in where and how workloads run. Every enterprise and every nation is waking up to the same reality. They need AI and cloud infrastructure they control, infrastructure no one can turn off or tamper with because of geopolitics.

During the quarter, we also announced strategic collaborations with NVIDIA, expanding our work across GPU-native analytics. In addition, we announced a strategic collaboration with Arm to expand how AI workloads run across IBM infrastructure. By enabling the Arm software ecosystem, within mission-critical environments like IBM Z, clients can scale AI closer to their data, while preserving the security and resilience they require.

These partnerships reflect our approach: open, flexible and on the infrastructure clients choose. We continue to make progress in Quantum and remain on track to deliver the first, large-scale fault-tolerant quantum computer by 2029. Here are some signposts of progress. In March, researchers used IBM Quantum hardware to simulate a 300 atom system with the Cleveland Clinic, demonstrating that Quantum computers can serve as reliable tools for pharmaceutical discovery.

Another team accurately simulated real magnetic materials. Magnetism is central to new forms of energy and electrification. These are significant demonstrations to-date that Quantum computers can serve as reliable tools for scientific discovery. We also released a new blueprint for quantum-centric supercomputing that outlines the architecture for integrating quantum and classical systems at scale. We strongly believe that our partners will achieve the first examples of quantum advantage this year, leveraging IBM hardware.

In closing, we are executing on our strategy of accelerating revenue growth and delivering higher profitability. Given our strong start to the year, we remain confident in our ability to sustain revenue growth of 5%-plus and grow free cash flow by about \$1 billion this year. With that, let me hand it over to Jim to go through the financials.

James Kavanaugh, IBM: Thanks, Arvind. In the first quarter, we delivered 6% revenue growth, 140 basis points of operating pre-tax margin expansion, 17% adjusted EBITDA growth, 19% diluted operating earnings per share growth, and \$2.2 billion of free cash flow, growing 13% year-to-year, representing our highest first quarter free cash flow in a decade and free cash flow margin in reported history. This performance reflects the work we have done to strengthen our software-led platforms, deliver innovation value to clients, and the durability of our financial model.

Now, I'll dive deeper into our segment performance. Software revenue grew 8%, marking a strong start to the year. This reflects the diversity of our portfolio, ongoing GenAI innovation, continued shift to higher growth end markets and flexible consumption model. Our ARR was solid at \$24.6 billion, up 10% since last year. Data revenue grew 16%, fueled by demand for our GenAI products, strengthen our strategic partnerships and inorganic contribution from DataStax and Confluent, which closed in mid-March. Red Hat growth accelerated 2 points sequentially to 10%, largely driven by the stabilization of consumption-based services revenue growth that we expected. OpenShift is now \$2 billion ARR business with strong growth. And virtualization continues to gain traction with over \$600

million of contracts signed since the beginning of 2024. Automation grew 7%, with February marking the one-year anniversary of the acquisition of HashiCorp. Over the last year, we have seen record HashiCorp bookings, leveraging IBM's go-to-market scale and achieved adjusted EBITDA accretion ahead of expectations. Transaction processing grew again, up 2%, as we monetize on the strong z17 program.

In Infrastructure, our revenue grew 12% this quarter, with Hybrid Infrastructure up 25% and infrastructure support down 6%. Within Hybrid Infrastructure, growth was broad-based with strong demand for our offerings across IBM Z, power and storage. IBM Z continues to outperform prior programs, growing 48% this quarter. Clients are investing in IBM Z as they modernize mission-critical workloads, driven by requirements for resiliency, security, and compliance, while enabling new AI capabilities on the platform. Distributed infrastructure grew double-digits with strength in both power and storage. Power growth was driven by demand for Power11, with its resiliency and performance advantages, supporting data-intensive workloads. In storage, growth reflected strong adoption of our new flash offerings introduced in the first quarter, which incorporate industry-leading agentic AI capabilities.

In Consulting, our revenue grew 1% this quarter, reflecting momentum in the business as client demand continues to shift towards enterprise-wide transformation. Signings returned to growth, up 6%, with strength across our application and data transformation offerings, driven by clients modernizing their environments to support AI adoption and capture value. Revenue growth was balanced across the portfolio, with both strategy and technology, and intelligent operations up 1%. Generative AI is now firmly integrated across our consulting engagements, representing about 30% of our backlog. This reflects how generative AI has become embedded in the work we do. Our differentiated asset-led delivery model continues to drive productivity and speed-to-value, combining deep domain expertise with software, automation, and reusable assets to help clients deploy AI securely and at scale.

Let me now discuss profitability. Several years ago, we set an ambitious objective to reinvent our enterprise operations for greater speed, lower friction, and structurally lower costs. Through disciplined execution, eliminating manual touch points, simplifying processes, and applying data automation and AI at scale, we have built a proven, repeatable, AI-enabled transformation engine that is accelerating. Since 2023, this has driven \$4.5 billion of productivity savings, with an additional \$1 billion expected in 2026. Our success is enabling us to accelerate investments in innovation, strengthen our competitive advantage as Client Zero and fuel our growth flywheel, while expanding our margins.

You can see this in the results this quarter, with productivity, revenue scale, and mix driving expansion of operating gross profit margin by 110 basis points, adjusted EBITDA margin by 170 basis points, and operating pre-tax margin by 140 basis points, all ahead of expectations. Segment profit margins expanded by 720 basis points in infrastructure and 60 basis points in software. Consulting segment

profit margin declined modestly, reflecting currency headwinds from geographic mix of the business and the reinvestment of productivity gains, amid an improving demand environment.

In the quarter, we generated \$2.2 billion of free cash flow, up about \$300 million year-over-year. The primary driver of this growth is adjusted EBITDA, up about \$600 million year-over-year, partially offset by higher net interest expense and increased investments in CapEx, as we expected coming into 2026.

We exited the first quarter, with a strong liquidity position and a solid investment-grade balance sheet with cash of \$11.8 billion. We invested \$10.5 billion in acquisitions, driven by the closing of Confluent and returned \$1.6 billion to shareholders in the form of dividends. Our debt balance ending the quarter was \$66.4 billion, including debt of \$12.8 billion for our financing business, with the receivables portfolio that is 80% investment-grade.

Let me now pivot to discuss our expectations going forward. The strong start to the year drives our confidence in delivering constant currency revenue growth of 5-plus-percent in 2026 and free cash flow growth of about \$1 billion year-over-year. Given where we are in the year, we believe it is prudent to maintain our guidance, even as the underlying performance and execution are off to an encouraging start. The combination of our focused portfolio, investment in innovation and our diversity across businesses drives the durability of our performance. Our revenue expectations are underpinned by our accelerating software business, which we now expect to grow 10-plus-percent this year.

In Consulting, the quality of our backlog and momentum in GenAI with backlog penetration at about 30%, continue to support an acceleration in revenue growth to low- to mid-single-digits for the year. We are off to a great start with z17. And four quarters into z17's launch, we prudently continue to expect infrastructure revenue to be down low-single-digits for the year, representing about 0.5 point impact to IBM. We remain confident this will be our strongest z cycle, given the AI innovation value we are delivering to clients.

The momentum in our productivity flywheel is fueling margin expansion, while enabling investment in innovation. Last quarter, we disclosed that we anticipated absorbing about \$600 million of dilution from Confluent in 2026, driven largely by stock-based compensation and interest expense. While we are absorbing incremental dilution, given the early closing of Confluent, actions we are taking to accelerate our cost synergies enable us to stay on track to expand operating pre-tax margins by about 1 point this year. Our operating tax rate for the year should be in the mid-teens, and the timing of discrete items can cause the rate to vary within the year.

For free cash flow, we continue to expect to grow about \$1 billion for the full year, driven primarily by growth in adjusted EBITDA. The headwinds I discussed heading into the year of higher cash taxes, higher CapEx, and higher net interest expense remain the same.

Looking to the second quarter, we expect our constant currency revenue growth rate to be similar to the full year. And for operating pre-tax margin, we expect about 50 basis points of expansion as

software mix and productivity are offset by dilution from the early closing of Confluent. Our second quarter operating tax rate should be in the mid-teens.

AI is fundamentally reshaping our clients' operating environments, increasing complexity, risk, and the need for flexibility. IBM's flywheel for growth built on trust, security and governance, a portfolio that helps enterprise put AI to work on their terms and sustained productivity that fuels rapid innovation, positions us to deliver value for our clients. We feel confident in our outlook and are excited about what's ahead. Arvind and I are now happy to take your questions. Olympia, let's get started.

Olympia McNerney, IBM: Thank you, Jim. Before we begin Q&A, I'd like to mention a couple of items. First, supplemental information is provided at the end of the presentation. And then, second, as always, I'd ask you to refrain from multi-part questions. Operator, let's please open it up for questions.

QUESTION AND ANSWER

Operator: Thank you. And at this time, we'll begin the question-and-answer session of the conference. [Operator instructions] And our first question comes from Amit Daryanani with Evercore ISI. Please state your question.

Amit Daryanani, Evercore ISI: Good. Thanks a lot. Good afternoon, everyone. Arvind, it's really nice to see the pickup in Red Hat growth and software acceleration broadly, especially as investors are debating software durability right now. When you step back and look at IBM's software portfolio, I'd love to understand, how would you characterize your mix between infrastructure versus applications, consumption versus subscription kind of stuff?

And as AI adoption really scales, where in that stack do you see the most incremental value accruing to IBM versus the ecosystem? If you just frame how you think of software, the puts and takes in a AI-centric world, it would be really helpful. Thank you.

Arvind Krishna, IBM: Good to hear from you, Amit. And as you can imagine, that's the question that occupies us, actually occupies our clients. And I know it occupies many investors' minds.

So let me first directly frame the answer in the dimensions that you laid out. If I think of infrastructure versus applications, I think if I count right, 4% of our portfolio, if I'm to be generous, could be called an application. Specifically, I think the only part of our portfolio that is applications would be Maximo.

And even that, which is looking at maintenance and asset management operations, even that many people would not really call an application, because it is the system of record as utilities and other people with very expensive infrastructure keep their maintenance records, including where a part may have come from 30 years ago.

So why do I say that? So if you look at the Red Hat portfolio, that's operating systems and container-based software and automation and runbook software, I think we would correctly call that, all-enabling software. The word 30 years ago would have been middleware, but that word has sort of gone away.

If I look at our data portfolio, it is databases, both relational and non-relational. And then, there is data movement. And then there is AI-enabling – by data movement, it's Confluent; and by AI, it's the watsonx portfolio. In automation, it's all about helping people take complexity out of how they manage their IT infrastructure, be that Turbonomic or Apptio or HashiCorp. And then, there is mainframe software, which is largely very similar to the first three I mentioned but for mainframe. If you look at these, that is all I think, in your words, I'll call it infrastructure, but I would call it more enabling software.

Second part of your question was on subscription versus consumption. I think our entire portfolio is very tied to consumption. We'll sometimes use the word capacity, because on mainframe, it is capacity, but capacity is equal to consumption, right? It's literally the MIPS that people use. It's not actually the installed capacity of the machine.

And off-the-mainframe, in the distributed world, whether on the cloud or on-premise, a lot of it is sometimes tied to amount of compute capacity that the software runs on, that's consumption in a large sense. Nobody's going to put it on processors, if you're not actually consuming it. And that's the vast majority of our software portfolio. So it is very much tied to that.

But then, I think the implicit question you're asking is, why would that get a tailwind from AI, as opposed to a headwind? So as people get serious about AI, because when they start experimenting, they may take a little bit of the data, they make a copy of it, they put it on a public cloud, they run it on some public frontier model, they get some results, and that's exciting to them.

As they get to scale, they've got to use the data from their internal systems. If they're using data from the internal systems, many parts of our portfolio, be it Red Hat, be it Confluent, will come up to be consumed more and more. As that gets consumed more and more, the automation part of the portfolio gets consumed more and more. As people do more and more fraud protection, not on sampling 1 in 10 transactions in the mainframe, but every single, that causes the mainframe consumption to go up and we can see that, by the way, in the mainframe numbers we printed in the first quarter.

So as we go through all of this, I think that this is a tailwind, because of the model that we picked. And by the way, I'll point out, this is not a model that's an accident of history. We have very consciously, over the last seven years, driven the portfolio into this, because we remain convinced that there is value in the underlying data layers, there is value in the business logic, and then, there's the interaction layer.

Value is going to decrease in that interaction layer, because as agents replace people for some fraction, we can debate how much of the interactions, then the interaction layer by itself is not sticky. The agents are going to be interacting much more with the underlying data and the business logic, and we sort of saw that coming six, seven years ago, and that is why we picked the portfolio we did.

I think that that hopefully gives you the sense and we can see it in our numbers that AI is structurally increasing the demand for the portfolio. But that is also why, both the organic products that we are building, for example, on software development and some of the acquisition targets we have had is to play into the tailwinds of AI demand.

Olympia McNeerney, IBM: Okay. Operator, let's take the next question.

Operator: Our next question comes from Wamsi Mohan with Bank of America. Please state your question.

Wamsi Mohan, Bank of America: Yes. Thank you so much. You cite greater than 10% growth in software now in 2026, the early close of Confluent itself should add about 1 point of growth just by itself. So how are you seeing the growth trajectory for the remainder of the software portfolio as we go through 2026?

And, Arvind, maybe quickly for you is IBM's appetite for M&A changing now that Confluent close is behind you, and given the broader de-rate in the software space?

James Kavanaugh, IBM: Thanks for the question, Wamsi. Appreciate it. Let's break down our software portfolio overall. First of all, we exit first quarter feeling very confident in our software portfolio, the innovation value, in the value proposition.

And I think that goes to the core of Amit's question and how Arvind answered it about how we're playing central to the thesis of where enterprise software and AI come together. That has always been predicated on our innovation strategy, our capital investment strategy, and our M&A strategy. And I think what you see in the first quarter is a reflection of the diversification of our portfolio and the durability of our software model coming out of the first quarter.

But if you dial back 90 days ago, what did Arvind and I say, entering 2026, we felt very confident about the strategic repositioning of software. Why? One, portfolio shift to higher growth end markets; two, strong annuity base that we've been building up, both organically and inorganically – by the way, we exit first quarter, approaching \$25 billion, growing 10%; three, new innovation and GenAI realization, and we could talk about that, M&A growth synergies. And now, we're into the second year of a very encouraging TP monetization opportunity.

But for the full year, how is it going to play out? We talked about entering the year 10% growth. Now, we see it growing 10-plus-percent accelerating.

Data. Data, we started out extremely strong, growing 16%. We are taking data up for the year. Yes, we closed Confluent early. Let's call it a couple of months overall. Closed it at the end of March, we were assuming somewhere in the mid-May time period – in the end of second quarter, we now see data up low-20-plus-percent range. That's going to deliver 5 points of software growth. That's representative of new innovation GenAI, the value of our platform-centric model and strategic partnerships, and then, also M&A contribution from Confluent, which would be about a little bit north of 15 points of that 20% to 25% growth overall for the year. So very strong organic.

Hybrid Cloud. Red Hat entered the year, we accelerated, delivered as expected. We posted 10% growth. Underneath that, contributed 2.5 points to IBM for the year, and we're well-positioned for that. Our subscription business accelerated. We got revenue under contract double digits. Red Hat OpenShift accelerating, \$2 billion ARR, virtualization now north of \$600 billion and consumption model returned back to expectation.

We are monitoring RHEL. RHEL did decelerate. I think that's a function of the federal lack of signings and the closure of the government in the fourth quarter that played through, but also a very dislocated hardware supply chain market. Automation on model, delivering over 2 points of growth. Hashi, great first year, record signings. We generated over \$200 million in new incremental ARR that should position 2026 well, new innovation, M&A growth synergies and NTP, continued growth, and we're off to a tremendous start, record start in our new z17. So we actually feel very good and more optimistic than where we were 90 days ago on software.

Arvind Krishna, IBM: Let me just address the M&A question, Wamsi, very quickly. Yes, the values that are out there right now are very attractive. That does not always mean that the sellers are willing to accept these values. That may take a few months for them to acknowledge that this is a new baseline. So if that's the case, I'm acknowledging that these are very attractive values.

Now, we have been a very disciplined acquirer. One, let us make sure that we fully integrate in and get all the benefits from Confluent. So that is going to take some months to get done. As we get through that and as the markets are at these values, that does open up our appetite perhaps more than it would in a normal year, but it's going to take a few months before we can go, acknowledge whether or not that's going to happen. And that's where I would sort of give you a bit of a color. So second half, if things stay where they are and if the values are where they are, maybe we can do something in the second half as we build up our cash balances, and we are 100% sure that Confluent is off to a strong start.

Olympia McNerney, IBM: Operator, let's take the next question.

Operator: Our next question comes from Ben Reitzes with Melius Research. Please state your question.

Ben Reitzes, Melius Research: Hey, guys. Thanks, I appreciate it. Arvind and Jim, I just want to talk about guidance. You guys rarely raise guidance after first quarter, I get it. I think there's just some concern out there as to are you seeing something in Europe that keeps you at bay right now? Are you seeing evidence of something slowing that keeps you from raising guidance? I mean, there's so many good things that are going on with regard to infrastructure and the software that you went through. So just wanted to clarify that.

And then, also with regard to guidance, the free cash flow was better than expected in the quarter. Why not raise it? Or do you just need to see more evidence? Thanks so much.

Arvind Krishna, IBM: Great. Ben, let me start out with just describing a little bit of the macro and what we are seeing as evidence and then why we are being a little bit prudent, and then, Jim will address all your questions on the specific guidance.

Let me start with the Middle East. We had the strongest growth we have seen in decades, not years, decades in the Middle East. So that gives you a sense that we are not seeing. There is no signal. I would tell you that I would expect the second quarter will play out similar to the first quarter in the Middle East. Our clients there, be it, the larger enterprises, be it government, they are clear they need to use and leverage technology to improve their own business.

In the first quarter, Europe was also strong. You can see that in the supplemental materials that we have provided. So there is nothing in the – what has already transpired. There has been no slowdown in deals. Deals have actually progressed at the rate and pace that we would want. If I look at pipeline and demand signals of the second quarter, we are not seeing any of this slowing down.

The only macro comment that we make is if the straits stay closed for another few weeks, then we know that there could be energy impacts in Europe, but that is speculative. That is not what we are seeing. And I expect that actually some of that we'll be able to absorb and maintain our acceleration, it's only if it crosses a certain level. So just based on only three months of the year have gone by is why we are making the prudent comment. Jim?

James Kavanaugh, IBM: Yeah. Ben, thanks for the question. Let's take a step back and put this in perspective, because I think you teed it up extremely well. And I've been in this role now nine years. Arvind has been in the role six, seven years. I don't think we've ever raised guidance in a first quarter.

But let's talk about the mentality, we've done a lot of work about strategically repositioning our portfolio, our business operating model, and the structural competitiveness of this business. And part of that was around how we were going to build discipline around execution in this company, and that execution mentality was around always a beat mentality at the end of the day.

The numbers speak for themselves in the first quarter, strongest first quarter revenue growth that we've had in over a decade. Arvind talked about the macroenvironment. Arguably, yes, we're operating in a dynamic world, and there's more uncertainty than there was 90 days ago, as we all know.

But within our lens of what we're looking at, we're executing extremely well across our high-value innovation software, infrastructure, and consulting that see signs of progress. Underneath that, look at what's happening to the fundamentals of our business, our operating margins are up 140 basis points. Our earnings are up nearly 20%, profits up 23%. This is an extremely strong start to the year.

And now, you get the free cash flow, which I know is a valuation measure as I spend a lot of time out with our investors talking about our strategic narrative and our financial investment thesis. Yes, free cash flow generation is the multiple that people more and more are valuing IBM and I would agree with that completely. We started out with the strongest free cash flow position in over a decade, highest free cash flow margin, up mid-teens.

Let's put this in perspective, less than 15% of what's required for the year. We got a lot of work ahead of us, but let's also put it in perspective, dial back a year ago, same call, same question. Look at how we executed on that mentality that Arvind's been trying to drive in this company. We had that same discussion. We executed well. We took off free cash flow throughout the year, and then, we blew through it in the fourth quarter.

And I will tell you coming out of the first quarter, there's no different mentality that we have here today. The underlying fundamentals, our adjusted EBITDA – by the way, all of this is high quality, sustainable, high-value realization overall. That is our free cash flow engine flywheel. That provides tremendous investment flexibility for us to continue to invest and drive long-term sustainable competitive advantage. And we don't see any different coming out of first quarter. But again, first quarter in, we're 90 days into an extremely important year, and our view is we should be prudent.

Olympia McNeerney, IBM: Operator, let's take the next question.

Operator: Your next question comes from Fatima Boolani with Citigroup. Please state your question.

Fatima Boolani, Citigroup: Good afternoon. Thank you for taking my questions. Arvind, I wanted to pull on a thread in your prepared remarks with respect to the mainframe potentially being a destination for more emerging use cases, especially around AI inferencing. So call them not your traditional or conventional mainframe use cases. I was hoping you could put some quantitative framing around that. What type of a workload mix are you seeing today that is you would consider conventionally mainframe, and what is the velocity of potential mix shift?

And then, as a related matter, as we think about the transaction processing and the MIPS growth momentum, how should that transpire and be expressed in the business in terms of the growth cadence for that particular segment as a follow through? Appreciate, there's a little bit of a lag there, but, would love your and Jim's comments on that. Thank you.

Arvind Krishna, IBM: Great. Thanks for the question, Fatima. Let me address the first part of your question, and then, I'll actually give it to Jim to address some of the quantification of those workloads.

So if we step back and look at it, over the last 60 years, mainframe has driven two great ways to monetize it. One has been what we call the classic MIPS or these are the compute parameters underneath that drive the transactional workloads that are great for the mainframe. Many people actually don't realize, but there are also, we call them Linux MIPS that are associated with the mainframe that people have been using to great effectiveness. But let me acknowledge, it is more Sparse Linux workloads as opposed to the highly, highly intense Linux workloads. AI is adding a third kind of compute capacity into the mainframe.

So just to make it very simple, today, if people are doing a payment authorization, almost all the credit card companies in the world use the mainframe for their credit card authorizations. If they want to do fraud, they can run a few rules in that engine, but then, they'll take a sampling of the transactions, let's

call it, 10% off the platform, because the latency that it introduces to take it off platform, you can't take them all. It'll just slow the whole system down. That's what they do off.

What happens if you could run a 20 billion, 30 billion parameter model right on the mainframe, suddenly because that is only milliseconds of latency, you can do that to every single transaction. So if you can take your fraud rate down from 50 basis points to 40 basis points, you can now do the math on what that is. They are all seeing that. So as I do that, we can do that for credit card authorizations. We can do it for retail banking transactions. We can do it for other payment operations. We can do it for claims and billing purposes.

So those are the workloads that are now coming on. So it is effectively a new capacity of the mainframe that previously was either very small, but outside the mainframe, or running on systems that are, what we would call, distributed infrastructure. We believe that this is going to play out. We see a large majority of our clients asking for the capacity. And currently, I believe we have a fully populated system, we can do about 450 billion inferences a day on the mainframe. So that gives you a sense of that.

We monetize that both through the extra hardware that is sold, but also by the supporting software for all of the AI inferencing that then runs on that increased capacity. So with that, then hopefully that gives you some color on what is happening, I'll give it to Jim.

James Kavanaugh, IBM: Yeah. Thank you, Fatima, for the question overall, I mean, mainframe modernization increases the strategic importance of IBM Z, as Arvind talked about. Why? Because the source of value is architectural, it's the platform. It's not the language. It's the tight integration of software, hardware, database, security, runtime, resiliency. And as Arvind talked about, this is a whole new modernization area of opportunity for us on that platform stack.

What is the driver of growth? Yes, 450 billion AI inferences, by the way, at one millisecond response time, 25 billion encryptions, transactions per day up to 8 nines of availability, quantum-safe encryption. And a TCO advantage of running it on mainframe on-prem versus the cloud anywhere from 3x to 15x depending on the size and complexity of that platform. So that's why, mainframe runs 73% of the world's transaction volumes in terms of value, 45 of the top 50 banks, 9 of the top 10 retailers, 4 to 5 top airlines, et cetera.

Now, you go to your second question about how do we monetize that value? One is the monetization of the platform of hardware. Arvind talked about AI MIPS. Second is that stack economic multiplier. Historically, we've been averaging about 3x to 4x stack multiplier for every hardware dollar we land.

Let me give you a stat. We just anniversaried our first full year of z17. That first full year is z17 versus the prior program z16 first full year, which by the way, was the best on record at that point in time. We've increased hardware placement value by over \$1 billion. Now, you take that \$1 billion and you think about the future monetization opportunity that we get. That's that 3x to 4x multiplier that will play out over time. A big chunk of that being our TP software, but it's also our storage attached. It's our

maintenance business, it's our financing business. We monetize that value based on how many MIPS are shipped in the market. And for four quarters in a row on z17, we've shipped over 100% growth of new MIPS in the market, including first quarter.

Why does that matter? Higher capacity is higher monetization opportunity, it's higher price opportunity, it's higher value creation opportunity. So we feel pretty good about that future monetization and multiplier effect as we play out 2026 and 2027.

Olympia McNeerney, IBM: Great. Next question, operator?

Operator: Our next question comes from Brent Thill with Jefferies. Please state your question.

Brent Thill, Jefferies: Hey, Jim, just on the constant currency for software, not to nitpick, but if you look at last year, 9% growth in Q1, I think it was 8% – 11% in Q4. So investors are asking, you're seeing a little bit of a downtick. Is that due to seasonality where maybe your contract signings were better in Q1, but maybe are being reflected in the reported numbers? Again, I know it's a modest deceleration, but anything to point out there?

James Kavanaugh, IBM: Yes, Brent, thank you very much for the question overall. I fully expected this one, because when you just look at the media print and the press release, fourth quarter, we posted a little over 11% growth. This quarter, we posted 8% growth. What gives? Do you feel still strength about your portfolio, your business, your investments, your new innovation? I think you nailed it right upfront.

One, understanding our business, our software portfolio, high-value recurring revenue, about 80% of our annual business, about \$30-billion-plus trailing 12 months, 80% of that is high-value annuity-based business. 20% is a transactional engine, underneath it. It's a big component of our perpetual license model, but it's a component of our subscription model, et cetera.

If you look at it, the entire three-point drop quarter-to-quarter is the fundamentals of the mix of the portfolio. In fourth quarter, we have about 30% of our business in the fourth quarter is transactional. In the first quarter, that's about 10%. When you look at the underpinnings of the core annuity by itself, we're actually accelerating that fourth quarter to first quarter. I think I said earlier on the call, our annuity ARR exiting first quarter approaching \$25 billion, that's up 10%, throughout the rest of the year, we'll go from a transactional quarter of about 10% first, and we'll peak probably in the fourth quarter of about 30-plus-percent.

On average, we'll be in the 20% overall. That will accelerate growth. That, coupled with M&A growth synergies, our GenAI portfolio, which has got a lot of momentum behind it, and our TP monetization and cycle, I would tell you, coming out of first quarter, I feel pretty good about 8% growth, and it positions us why we said 90 days ago, confident in 10%. Now we're saying, yes, we closed Confluent earlier, and we're confident now in accelerating that to 10%-plus.

Olympia McNeerney, IBM: Operator, let's take the next question.

Operator: Our next question comes from Erik Woodring with Morgan Stanley. Please state your question.

Erik Woodring, Morgan Stanley: Awesome. Thank you for taking my question tonight, guys. Jim, you briefly alluded to it earlier, but can you maybe just detail how IBM is broadly managing and/or mitigating some of these supply chain headwinds, whether that's higher memory costs or supply challenges? Meaning, how material is memory within the infrastructure biz? How are you mitigating? How are customers responding? How does it impact your outlook on growth and margins? If you could just maybe dig into us, that would be super helpful. Thank you very much.

James Kavanaugh, IBM: Yeah, absolutely. Thank you, Erik, for the question overall. You understand our business extremely well, underneath Arvind's leadership, we have strategically repositioned this portfolio. It's been a lot of work around portfolio optimization – by the way, that's both leveraging the strength of our cash flow, our financial flexibility to buy high-value, innovative-based companies in category-leading technologies with structural growth profiles that help IBM, but it's also around divestitures of portfolio.

But where I'm going with this is today, when you look at IBM's portfolio, we're a human capital asset, IP-based business, that's 75% on a way to 80%, by the way, underneath that software, 45% on its way to 50%-plus overall. Our hardware business is extremely important as a value creator to IBM, but top line, it's about 25% of our business. But that's high-value innovation on mainframe platform overall.

Now, you look underneath that around the supply chain dislocation, around commodity cost increases in particular around memory, it has a de minimis impact to us overall. Think about mainframe overall. Well, it impacts storage and potentially some components of our distributed infrastructure. Absolutely. But look underneath it, we're able to, one, we've been in existence for 115, 116 years overall. We know how to run global supply chains. We drive supplier optimization, supply chain diversification, procurement strategies overall. And I think we've been able to mitigate this dislocation overall.

The area we're watching it is in the software area around RHEL. I mean, RHEL is tied to enterprise hardware placements overall. And we'll continue monitoring that. But look at our hardware performance, we accelerated growth at 15%, our distributed infrastructure, at actual rates grown 17%, constant currency grown 13%. By the way, I didn't even talk about it, our infrastructure pre-tax margins are up 720 basis points. So we know how to manage global supply chains and commodity costs inside the company and extract value overall.

Arvind Krishna, IBM: Jim, let me just add a couple of sentences to your statement. Erik, Jim mentioned that we have worked with a lot of these suppliers for a number of years and decades. They like working with us partly because of the relationships we have built up with them over the years, but also, because we help them stress-test new capabilities. And they like the fact that our systems are very high-performing, because that gives them brand reputation as they go out to the wider market.

That does help, not completely, but somewhat mitigate some of the supply chain constraints, because we are early users of their newest memory technologies.

Olympia McNerney, IBM: Operator, next question.

Operator: Your next question comes from Jim Schneider with Goldman Sachs. Please state your question.

Jim Schneider, Goldman Sachs: Good afternoon. Thanks for taking my question. I was wondering if you could maybe comment on the AI bookings, which is a metric you previously given – you've given, but I think you just commented as a percentage of your total bookings right now, did that accelerate or decelerate in the quarter?

And then, maybe just kind of comment on any update you see for the consulting business this year, either given more macro uncertainty, do you expect any kind of diminution in the growth rate you expect this year? Thank you.

James Kavanaugh, IBM: Yeah. Thanks, Jim. As we talked about, we exited last year with the book of business around AI, which as you know, we talked about consulting and software within that vernacular. I think it was important over the last couple of years, because as the explosion of GenAI hit, we had to give a perspective about whether we were winning and capitalizing on participating in that market.

We exited last year, what, \$12.5 billion over \$12.5 billion book of business. But now let's bring it back because in January, we talked about, it's embedded across our portfolio. It's embedded in software. It's central thesis to how we run our consulting business right now. It's embedded across our infrastructure business. And we said coming into 2026, we were going to talk about it more from an outcome base, revenue base, and value contribution base overall.

But let's talk about software. Software – GenAI continues to be a tailwind overall. The positioning of our portfolio with the explosion of AI, with the applications agents, with us owning the foundational layer of Linux containerization, you see that play out with the acceleration of our Red Hat OpenShift business now \$2 billion, growing north, I think, high-20% growth overall. Second, the importance of the data layer, Arvind talked about Confluent positioning us to be the cross-platform as a data connector, automation, the need for resiliency, observability, FinOps.

Software, let's talk about, one, it's accelerating our growth profile overall, but let me put some numbers behind it versus just a overarching book of business. Our software book, from an annualized revenue trailing 12 months, we finished last year at \$30 billion, right? 80% of that, as I said earlier, high value recurring revenue, 20% transactional. Read that about \$6 billion.

Over the last trailing 12 months on an accelerating basis, our AI platform, agents, assistants, orchestration is north of \$1.5 billion. It's already about 25% penetrated in our software business, growing north of 40%. It's contributing 2 points of growth on an annualized basis. And the thing we love about it, it has a multiplier effect over time, so it's an acceleration there.

Consulting. Consulting is about 40% of our signings. 30% of our backlog is GenAI now, over 20% of our revenue. And on a ARR revenue perspective, in the first quarter, we eclipsed \$4 billion ARR. So it is central to the way we run a services-as-a-software model overall.

And in Infrastructure, both Arvind and I talked about, it's embedded on the chip of z17, Spyre inferencing. I think Arvind talked about it in prepared remarks. Clients that have implemented watsonx Code Assistant for Z, were seeing 3x differential on growth and capacity. And you see in our distributed infrastructure, we're accelerating growth.

Now, your second question around Consulting. We are seeing signposts of progress overall. One, our demand profile, our backlog quality, our GenAI, which I just talked about, our strategic partnership headroom opportunity, our portfolio mix composition more at a higher growth areas, and our services-as-a-software model, which we think we have an industry-leading position with our IBM Consulting Advantage platform.

Well, let me put some stats on it. One, signings, we returned to growth. Great quarter overall, large transformational deals around GenAI. The health and mix of net new business and expansions up 7 points year-to-year, up 4 points quarter to quarter. 400 new clients captured in the first quarter. Our backlog quality overall, our erosion is stable. Our duration continues to come down. Our backlog realization is actually accelerating throughout the year. And our backlog yields are up 4 points year-over-year, talking about the quality and value we're able to deliver.

And I talked about GenAI, 80% of our GenAI book of business right now is coming from captured from net new clients overall. And I'll stress that over \$4 billion revenue ARR. So that positions our confidence in the year of us accelerating our revenue growth, around low-single-digits. And if things go well, can we do better than that? Obviously, yes.

Olympia McNerney, IBM: Great. Operator, let's take one last question.

Operator: Our last question comes from Matt Swanson with RBC Capital Markets. Please state your question.

Matt Swanson, RBC Capital Markets: Great. Thank you so much for squeezing me in here. Arvind, it's really interesting going over the software segments, when you showed how low of an exposure you have to the application space. There's obviously been a ton of debate right now around who's going to kind of win the GenAI workloads of applications.

We've seen you operate at such a strong kind of Switzerland foundational player in the hybrid cloud. When we look at AI, like, how are you setting IBM up to win? Kind of regardless what ends up being the winner of the GenAI application layer? And what kind of investments does that take?

Arvind Krishna, IBM: Matt, thanks for that question. So we made the decision about three years ago that we were going to be neutral and Switzerland-like also on our usage of frontier models. Because I think when you are saying that GenAI applications, I think, for many people that is synonymous with

the frontier model providers, not just the frontier models, but all the surrounding software accoutrements that all of them are giving.

So we are going to play where clients want to be hybrid, where clients may want to function across multiple clouds or also because of either sovereignty, or brand, or privacy, or in the end, economics. They might also have a private addition – in addition to what they use on public. So as we go across that, we are building, for example, our software development AI product, Project Bob. It is out. We actually chose not to announce it. Nevertheless, 200 people signed up to use it. So that gives us a signal that we have something.

Now, why would they use us as opposed to just one of the Codex or equivalent models, is if they also have a lot of code that they do not want to actually take out in public, and also, they want to address the entire software development – meaning, including testing, including patching, including documentation, including maintenance are the kinds of things that we provide. Ditto, as we look at how they might want to use agents that come inside their enterprise, then we use Confluent to go manage and control how they expose data from inside things.

So as we sort of look at that, Matt, I think we are very clear-eyed. There will be people who will be frontier model providers. You can debate how those half a dozen or a dozen; today, it's somewhere in that range. We actually do not want to even predict which of them will be the eventual winners. We want to work with all of them.

Then, we also work with open-weight models. And we produce models where we have either domain expertise or people may want much smaller models to be able to run them on-premise, or I'll say, euphemistically, on a 1-to-4 GPU server node as opposed to a very, very large cluster. So that's the model place. We are going to then help our clients deploy these models to gain value.

As we have unlocked, Jim talked about the \$4.5 billion of internal value, how do you reduce your total tax expense? How do you reduce procurement expense? How do you reduce accounts payable? How do you reduce quote-to-cash? As we walk across these processes, we get a lot of knowledge on how to capture that into agents, but then, we are not going to be fixated. Whichever model you want to use, you can use. And wherever you want to run them, we'll help you run them. And we think that's a good half of the world is interested in that paradigm. And that's how, Matt, we are going to be able to go win in this world as it unfolds going forward.

So just to close, the innovation value we are delivering to our clients and our strong start to the year reinforce our confidence in our growth trajectory. We look forward to continuing this dialogue as we move through the year.

Olympia McNerney, IBM: Thank you, Arvind. Operator, let me turn it back to you to close out the call.

Operator: Thank you for participating on today's call. The conference has now ended. You may disconnect at this time.